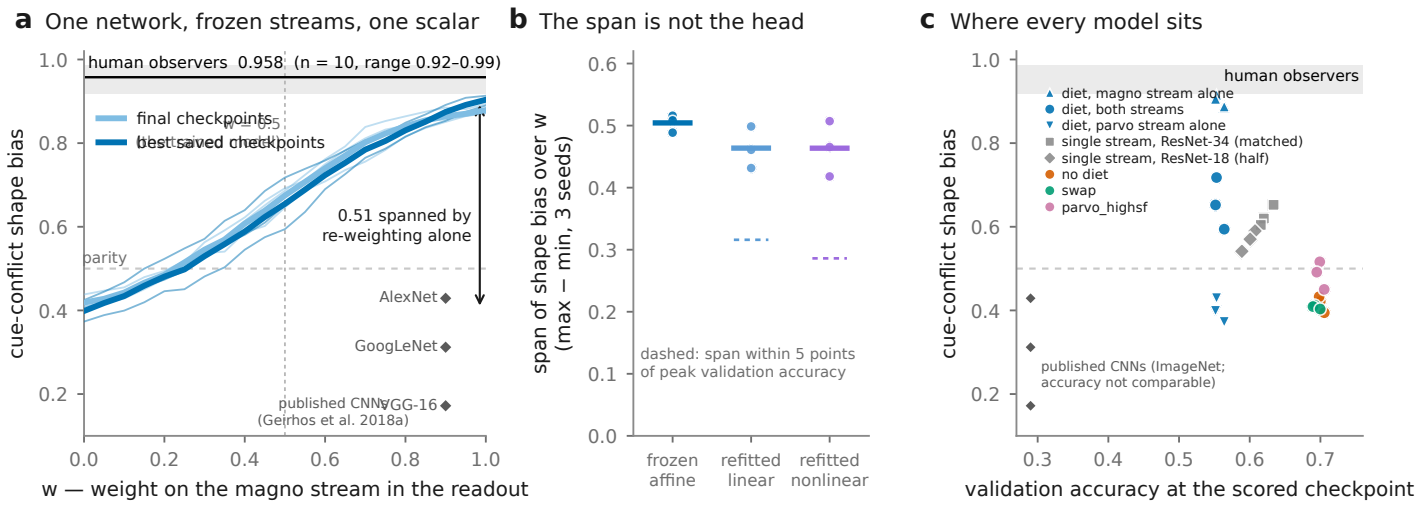
a One network, frozen streams, one scalar
human observers 0.958 (n = 10, range 0.92–0.99)
final checkpoints
best saved checkpoints
w = 0.5
0.51 spanned by
re-weighting alone
parity
AlexNet
GoogLeNet
VGG-16
published CNNs
(Geirhos et al. 2018a)
cue-conflict shape bias
w — weight on the magno stream in the readout
b The span is not the head
span of shape bias over w
(max − min, 3 seeds)
dashed: span within 5 points
of peak validation accuracy
frozen
affine
refitted
linear
refitted
nonlinear
c Where every model sits
human observers
diet, magno stream alone
diet, both streams
diet, parvo stream alone
single stream, ResNet-34 (matched)
single stream, ResNet-18 (half)
no diet
swap
parvo_highsf
published CNNs (ImageNet;
accuracy not comparable)
cue-conflict shape bias
validation accuracy at the scored checkpoint