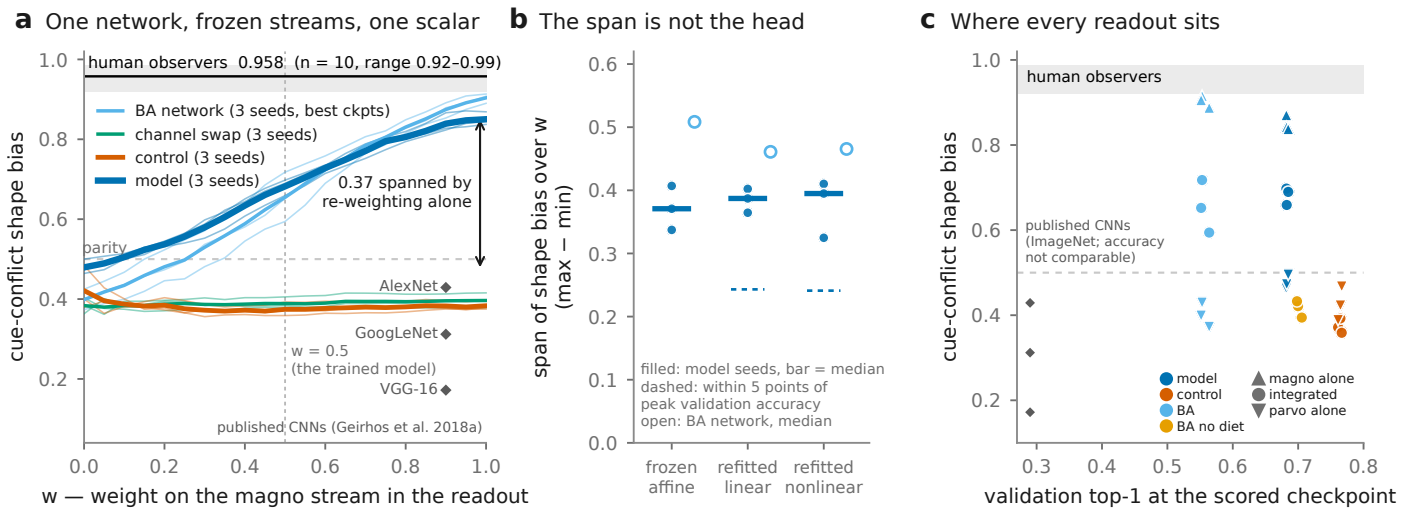
a One network, frozen streams, one scalar
human observers 0.958 (n = 10, range 0.92–0.99)
BA network (3 seeds, best ckpts)
channel swap (3 seeds)
control (3 seeds)
model (3 seeds)
0.37 spanned by
re-weighting alone
parity
AlexNet
GoogLeNet
w = 0.5
(the trained model)
VGG-16
published CNNs (Geirhos et al. 2018a)
cue-conflict shape bias
w — weight on the magno stream in the readout
b The span is not the head
span of shape bias over w
(max − min)
filled: model seeds, bar = median
dashed: within 5 points of
peak validation accuracy
open: BA network, median
frozen
affine
refitted
linear
refitted
nonlinear
c Where every readout sits
human observers
published CNNs
(ImageNet; accuracy
not comparable)
model
control
BA
BA no diet
magno alone
integrated
parvo alone
cue-conflict shape bias
validation top-1 at the scored checkpoint