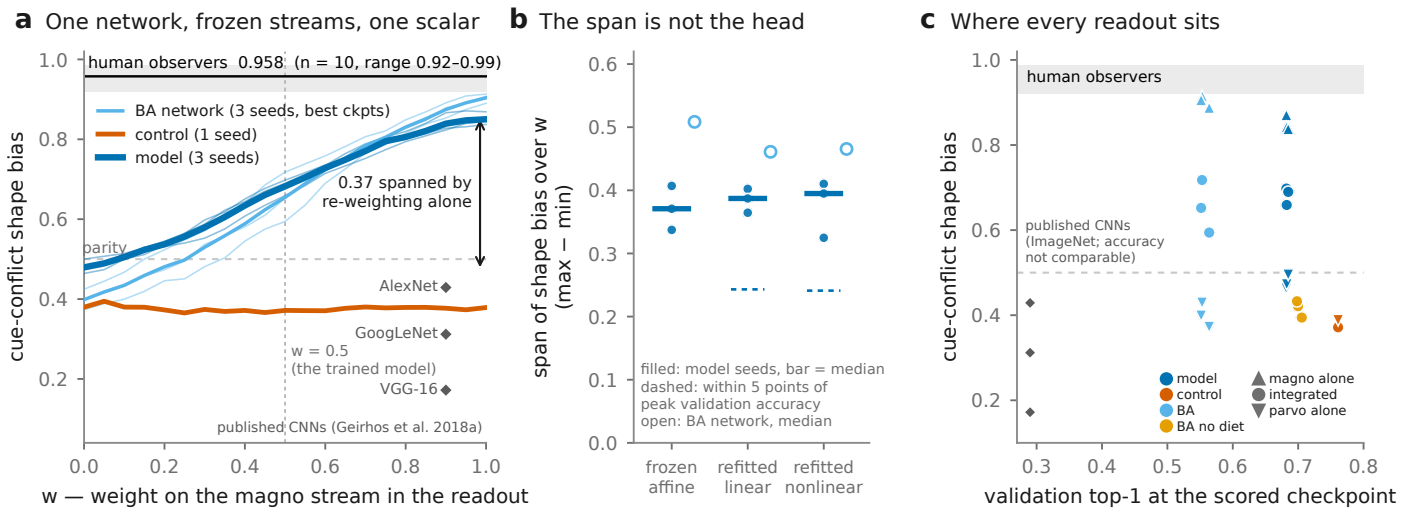
a One network, frozen streams, one scalar
human observers 0.958 (n = 10, range 0.92–0.99)
BA network (3 seeds, best ckpts)
control (1 seed)
model (3 seeds)
0.37 spanned by re-weighting alone
parity
AlexNet
GoogLeNet
w = 0.5 (the trained model)
VGG-16
published CNNs (Geirhos et al. 2018a)
cue-conflict shape bias
w — weight on the magno stream in the readout
b The span is not the head
span of shape bias over w (max − min)
filled: model seeds, bar = median
dashed: within 5 points of peak validation accuracy
open: BA network, median
frozen affine
refitted linear
refitted nonlinear
c Where every readout sits
human observers
published CNNs (ImageNet; accuracy not comparable)
cue-conflict shape bias
validation top-1 at the scored checkpoint
model
control
BA
BA no diet
magno alone
integrated
parvo alone